

Assumptions:

- i). There must be linear relationship btwn independent and dependent variables
- ii). Mean of residual must be 0 or close to 0
- iii). Variance around the regression line is the same for all values of X.
- iv). There should be no multicollinearity

Multiple Linear Regression

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: df=pd.read_csv("50_Startups.csv")
```

```
In [3]: len(df)
```

```
Out[3]: 50
```

```
In [4]: df.head()
```

```
Out[4]:
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

```
In [5]: df.shape
```

```
Out[5]: (50, 5)
```

Importing required Library

```
In [6]: from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error


import matplotlib.pyplot as plt
```

Visualizing Data


```
In [7]: plt.scatter(df['Marketing Spend'],df['Profit'],alpha=0.8,color='red')
plt.title("Marketing Spend vs Profit")
plt.xlabel("Marketing Spend")
plt.ylabel("Profit")
plt.show()
```



```
In [8]: plt.scatter(df['Administration'],df['Profit'],alpha=0.8,color='red')
plt.title("Admininstration vs Profit")
plt.xlabel("Administration")
plt.ylabel("Profit")
plt.show()
```




```
In [9]: plt.scatter(df['R&D Spend'],df['Profit'],alpha=0.8,color='red')
plt.title("R&D Spend vs Profit")
plt.xlabel("R&D Spend")
plt.ylabel("Profit")
plt.show()
```



```
In [10]: ax=df.groupby(['State'])['Profit'].mean()
ax.plot.bar(figsize=(9,5),fontsize=14)
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x2102b1d6b00>
```



```
In [11]: df.State.value_counts()
```

```
Out[11]: New York      17
California    17
Florida       16
Name: State, dtype: int64
```

```
In [12]: #Creating Dummy Variable
df['NewYork']=np.where(df['State']=="New York",1,0)
df['California']=np.where(df['State']=="California",1,0)
df["Florida"]=np.where(df["State"]=="Florida",1,0)

#Drop State
df.drop(columns=["State"],inplace=True)
```

```
In [13]: dep="Profit"
ind=df.columns.tolist()

#Removing Profit from independent variable
ind.remove(dep)
print(ind)
```

```
['R&D Spend', 'Administration', 'Marketing Spend', 'NewYork', 'California', 'Florida']
```

Splitting the Data

```
In [14]: x=df[ind]
y=df[dep]
```

```
In [15]: xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.2,random_state=0)
```

```
In [16]: #Transforming Data (helps to better fit the model)
scaler=MinMaxScaler()
xtrain=scaler.fit_transform(xtrain)
xtest=scaler.transform(xtest)
```

C:\Users\Gaurav\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:334:
DataConversionWarning: Data with input dtype int32, float64 were all converted
to float64 by MinMaxScaler.
return self.partial_fit(X, y)

Building the Model

```
In [17]: regressor=LinearRegression()
regressor.fit(xtrain,ytrain)
```

```
Out[17]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```

Predicting on Test Data

```
In [18]: pred=regressor.predict(xtest)
```

```
In [19]: from sklearn.metrics import mean_squared_error
print("    Root Mean Squared Error")
print("-----")
np.sqrt(mean_squared_error(pred,ytest))
```

```
    Root Mean Squared Error
-----
```

```
Out[19]: 9137.990152794935
```

#On an average we will be off by 9k

```
In [20]: print("          R2 Score")  
print("-----")  
np.sqrt(r2_score(pred,ytest))
```

R2 Score

```
Out[20]: 0.9640409332242128
```

R2 will be 1 for perfect model.It is close to 1.

Look for Gradient Descent Algorithm for knowing how the regressor is calculated internally